

FREE / 5-LINE WORKSHEET

Does this agent decision need a model?

Use this on one real routing decision. Tap each blue line to type a short answer, or print the page to write by hand. No account, email, or Jev key is needed.

Fill in the five lines

1

Input

What trusted data and ambiguous text does your application receive?

2

Allowed outputs

Which stable IDs may it return? Include a review route when needed.

3

Exact rules

Which cases can your code settle from trusted data?

4

Failure route

What happens on an unknown ID, low confidence, timeout, or error?

5

Final authority

Which code validates arguments, permissions, and resources before action?

If you cannot name the allowed outputs or failure route, narrow the task before choosing a model.

Choose the simplest fitting method

Decision shape	Start with
Exact rule over trusted data	Deterministic code
Ambiguous meaning; finite outcomes	Evaluate a bounded decision model
Writing, synthesis, open reasoning	Generative model
High-impact action	Deterministic policy + approval

A model's selected label or tool never grants permission to act.

WORKED EXAMPLE / SUPPORT ROUTING

One ticket. Four queues.

Illustrative ticket: "I was charged twice after the app froze."

Input

Ticket text, plus invoice or account fields the application has already verified.

Allowed outputs

billing / technical / account / review

Exact rule and semantic case

If a verified duplicate-charge event and your policy settle the route, branch in code. If the relevant meaning remains ambiguous, evaluate a bounded choice over the four IDs. Treat its answer as a proposal.

Failure and authority

Unknown output or uncertainty goes to review. A timeout follows a bounded retry or review policy. The application validates the ID and arguments, checks the user's access to the account, and only then routes or asks for review.

Compare before replacing your current path

Label representative cases before reading model results. Run the same cases through the current route and candidate. Record one row per case:

case_id, expected_route, selected_route, valid_output, correct_route,
fallback, end_to_end_ms, cost_basis

Report the sample size and definition of every rate. Include retries, review, and provider time in the full path. Keep tuning cases separate from evaluation cases.

Inspect a free runnable example

The public [TypeScript support-ticket router](#) uses synthetic input, a local allowlist, and human review. It prints a recommendation and performs no action.

Optional next step

For TypeScript and Python quickstarts, 16 decision recipes, fallbacks, and an offline harness for your own measurements, see the paid [Jev Operator Kit](#) (US\$19 one-time). The worksheet and kit share an author. Jev access and provider billing are separate. No Jev benchmark result or future update is promised.

This ticket is illustrative, not a recorded Jev decision or benchmark result.